

## Problem Statement

You are tasked with building an offline-capable, local-first AI application—KCC Query Assistant—that allows users to query agricultural advice from the Kisan Call Center (KCC) dataset.

## Objectives

1. Data Integration & Preprocessing
  - Download the public KCC dataset.
  - Clean, normalize, and split into logical “document” chunks or Q&A pairs.
  - Export both raw and preprocessed formats, preserving metadata fields.
2. Local LLM Deployment
  - Use an open-source model via the Ollama API (e.g., Gemma 3, Deepseek).
  - Quantize as needed for CPU/GPU efficiency (e.g., GGUF/GPTQ).
  - Ensure the model runs entirely offline.
3. Retrieval-Augmented Generation (RAG)
  - Generate embeddings for each document chunk using a sentence-transformer (e.g., MPNet, bge-large-en).
  - Store embeddings in a lightweight vector database (ChromaDB, FAISS, or MongoDB).
  - Implement semantic search to retrieve top-k relevant chunks for any query.
  - If no context meets a relevance threshold, invoke a live Internet search and clearly notify the user.
4. User Interface
  - Build a simple local web app (Streamlit or React.js).
  - Allow natural-language queries.
  - Display structured answers, highlighting which advice came from KCC versus fallback search.

## Flow Summary

1. Data Ingestion
  - Load raw KCC CSV

2. Preprocessing  
→ Clean, normalize, chunk Q&A pairs
3. Embedding Generation  
→ Encode chunks into vectors
4. Vector Store Ingestion  
→ Index embeddings in FAISS/ChromaDB/MongoDB
5. Query Handling (UI Layer)
  - User submits query
  - Retrieve top-k context via semantic search
  - If context found:
    - Pass context + query to LLM → display response
  - Else:
    - Notify “No local context found”
    - Perform live Internet search → display fallback results

## Example Use Cases

- “What pest-control methods are recommended for paddy in Tamil Nadu?”
- “How to manage drought stress in groundnut cultivation?”
- “What issues do sugarcane farmers in Maharashtra commonly face?”

You may create additional queries reflective of your processed data. Ensure your UI, documentation, and demonstrations cover these core scenarios.



## Submission Guidelines

Please follow these instructions carefully when submitting your KCC Query Assistant project.

### 1. Repository Setup

- Create a private Git repository named `YourName_KCCQueryAssistant`.
- Grant our review team collaborator access.
- Submit only the repo link and your Google Drive demo-video link via the form.

### 2. Deliverable Checklist

Ensure your repository contains:

- Raw KCC Dataset plus comprehensive technical documentation covering every step workflow.
- Preprocessed Dataset files and scripts.
- Vector Database Artifacts or build scripts (ChromaDB, FAISS, or MongoDB).
- Source Code for:
  1. Data ingestion & preprocessing
  2. Embedding generation & vector store ingestion
  3. RAG pipeline & LLM integration
  4. Web UI (Streamlit or React.js)
- Sample Queries relevant to data to test atleast 10.
- README.md covering installation, dependencies, overview, and launch instructions.

### 3. Demonstration Video

- Record a 3–5 min screencast showing:
  1. Local startup of the LLM and vector store
  2. At Least 3-5 Queries returning KCC-based answers
  3. Fallback to live Internet search when no local context is found for at least 2-3 queries.
- Upload to Google Drive (view-only) and include the link.